



Level 5-3

# Gophers & Friends

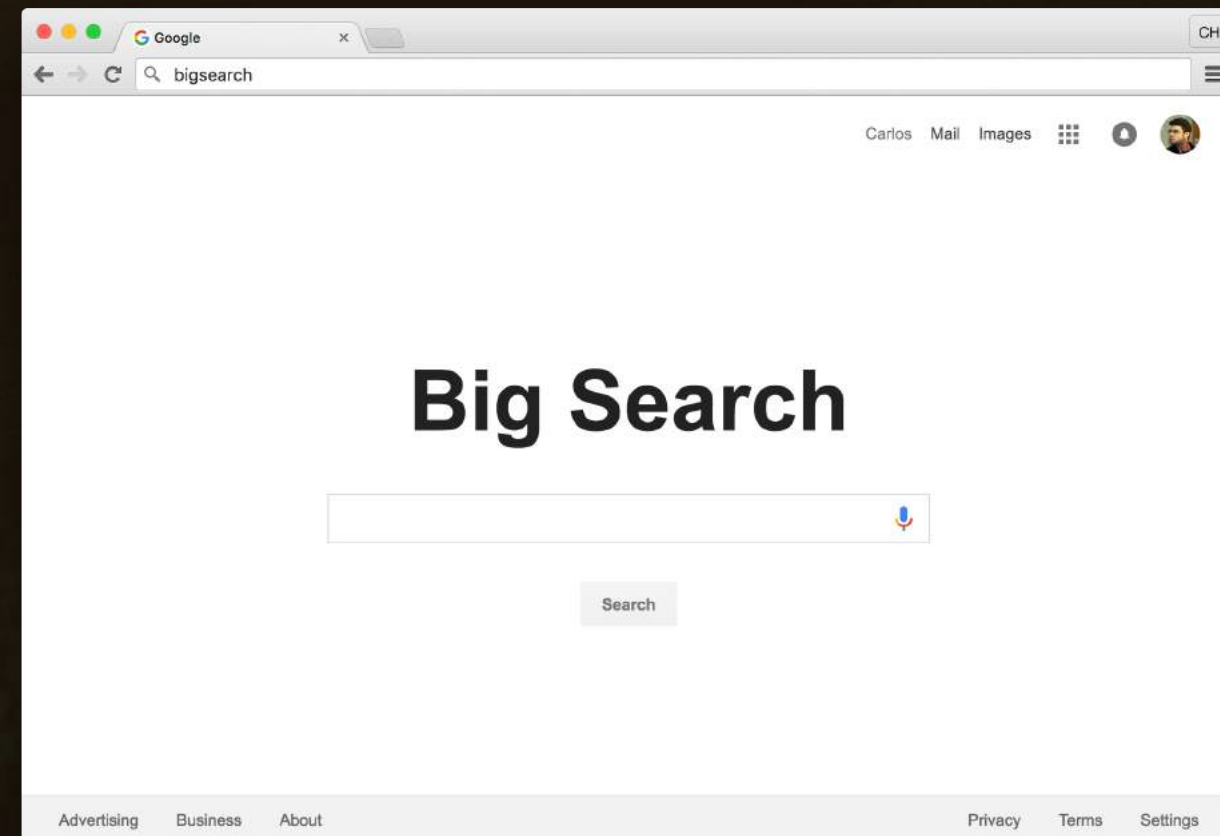
goroutines

ON TRACK  
with  
GOLANG

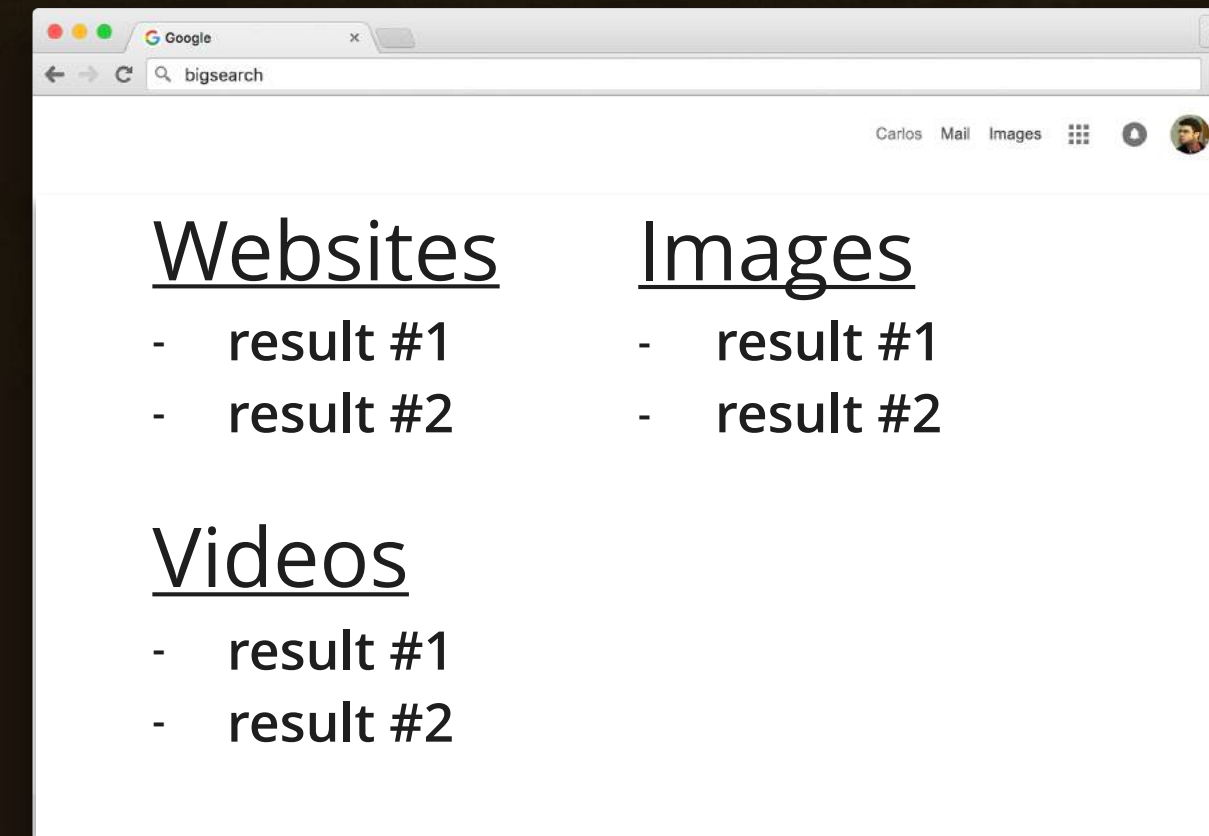


# The “Big Search” Website

This search engine will search the web for websites, images, and videos.



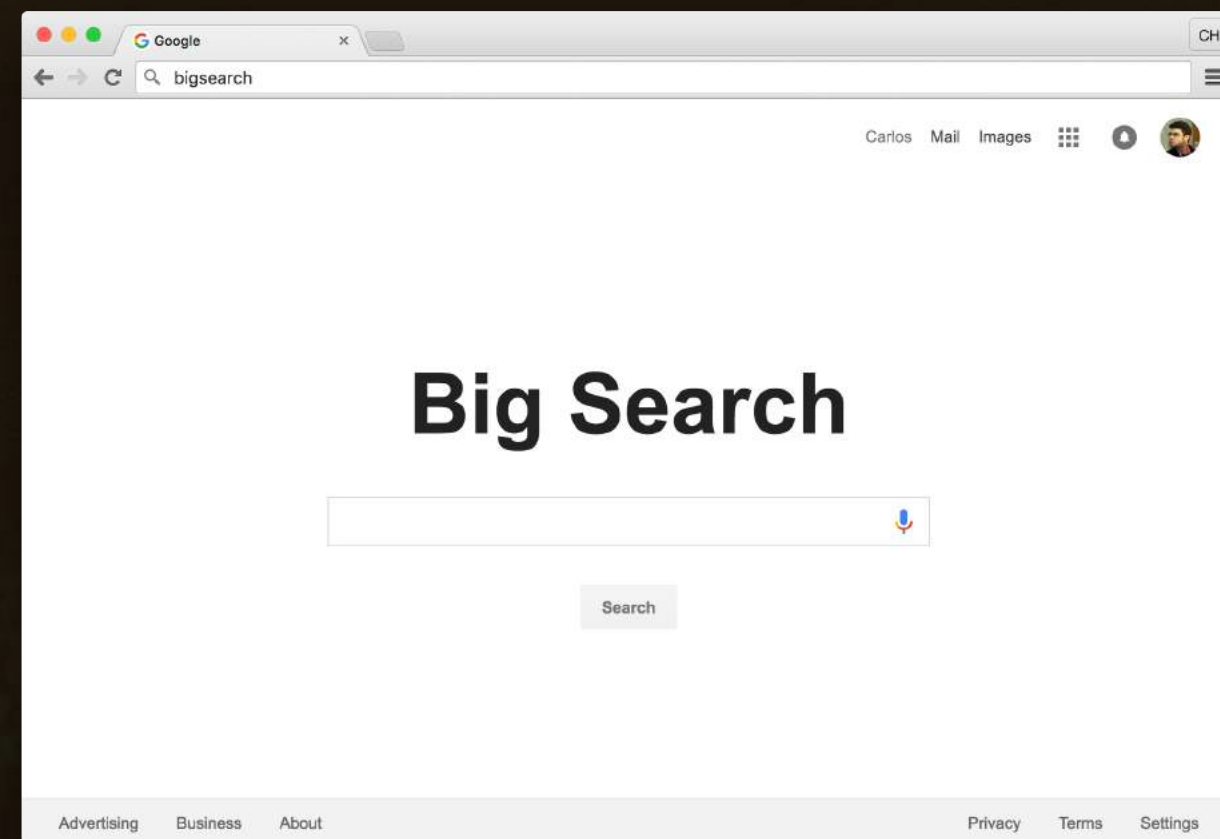
*Display results for a search*



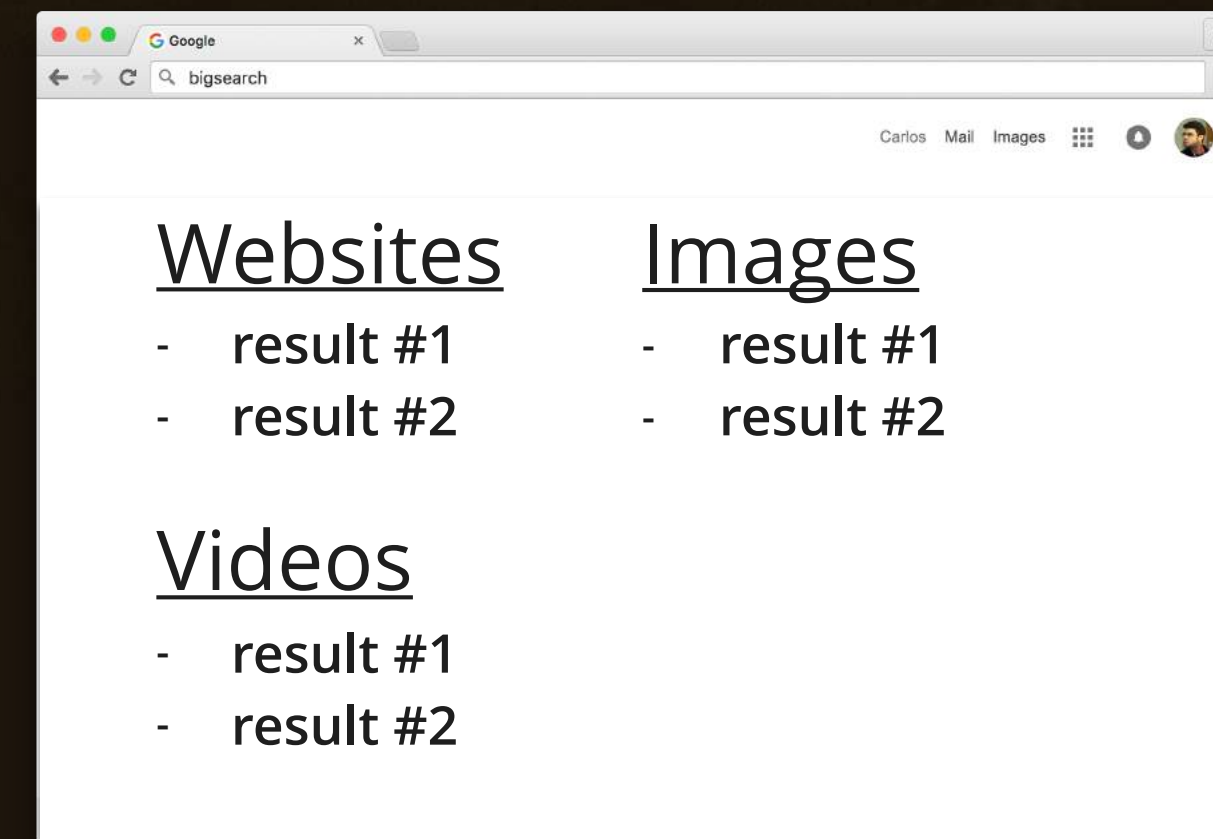


# Sequential Programs


In sequential programs, before a new task starts, the previous one **must finish**.



*Searches websites, images,  
and videos one at a time*



*(Total: 9 seconds)*

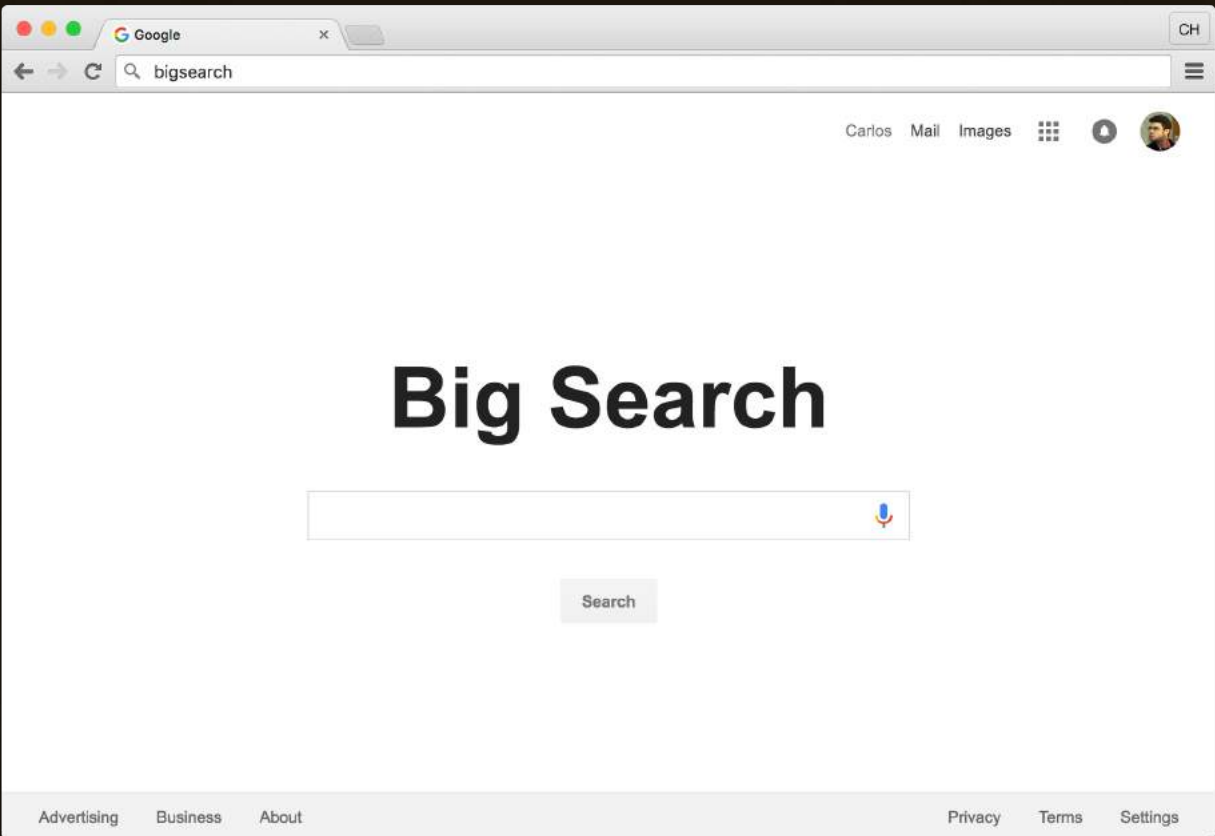


*Takes as long as the  
sum of all of tasks*

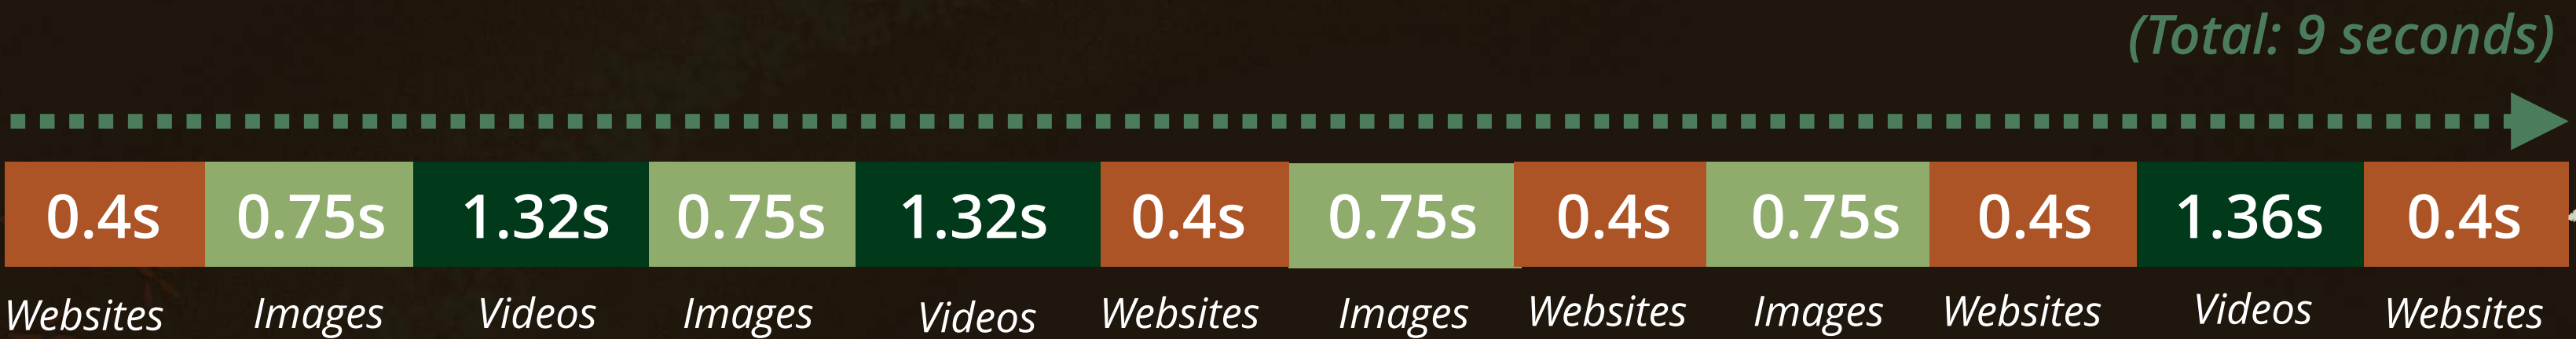
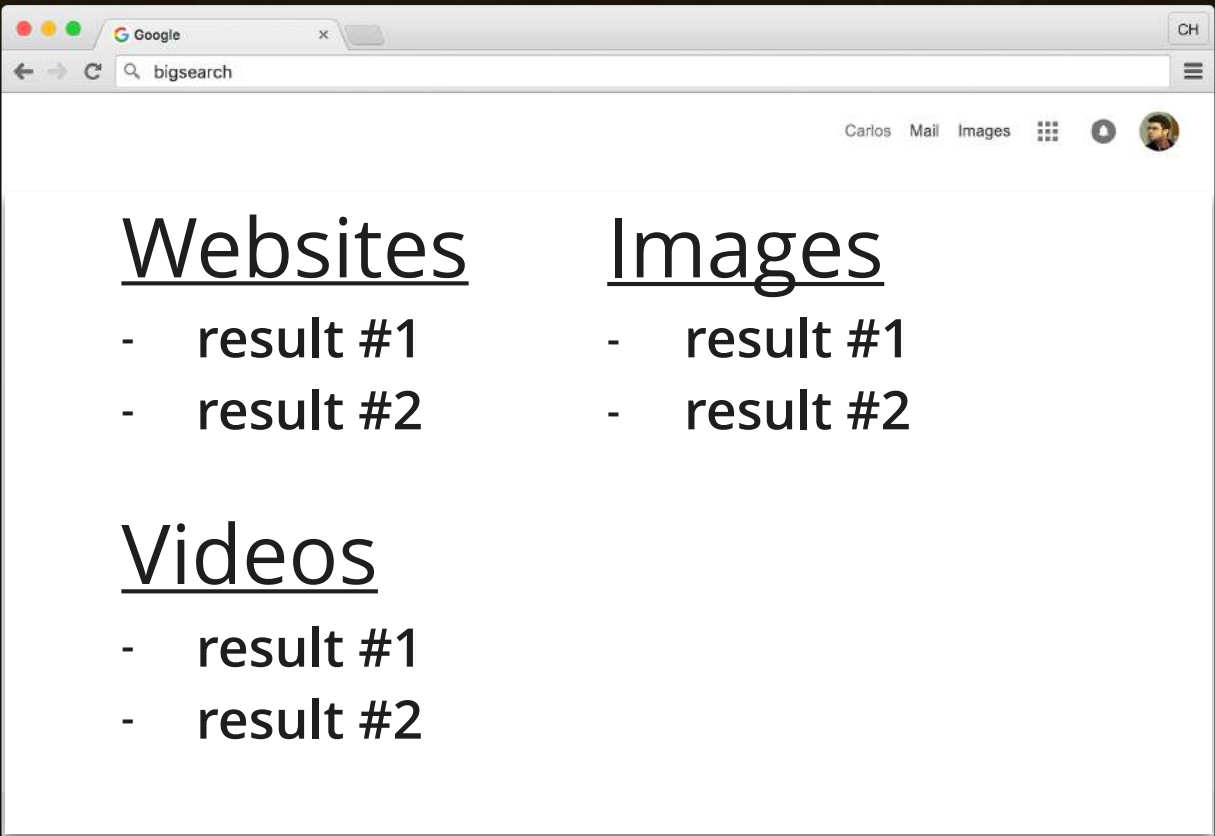


# Concurrent Programs

In concurrent programs, multiple tasks can be **executed independently** and may **appear simultaneous**.



*Searches websites, images, and videos "at the same time"*



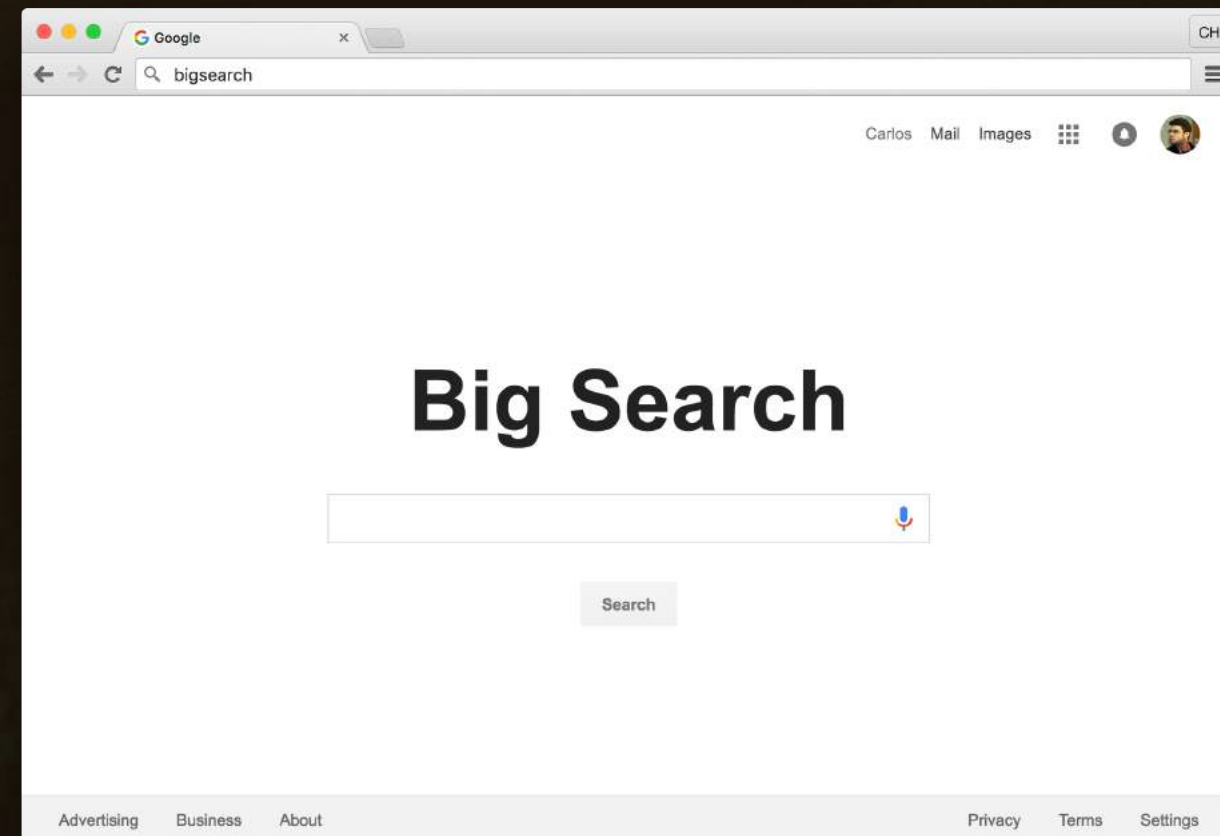
*Can be executed in any order*

*Also takes as long as the sum of all tasks, but tasks alternate time slices*


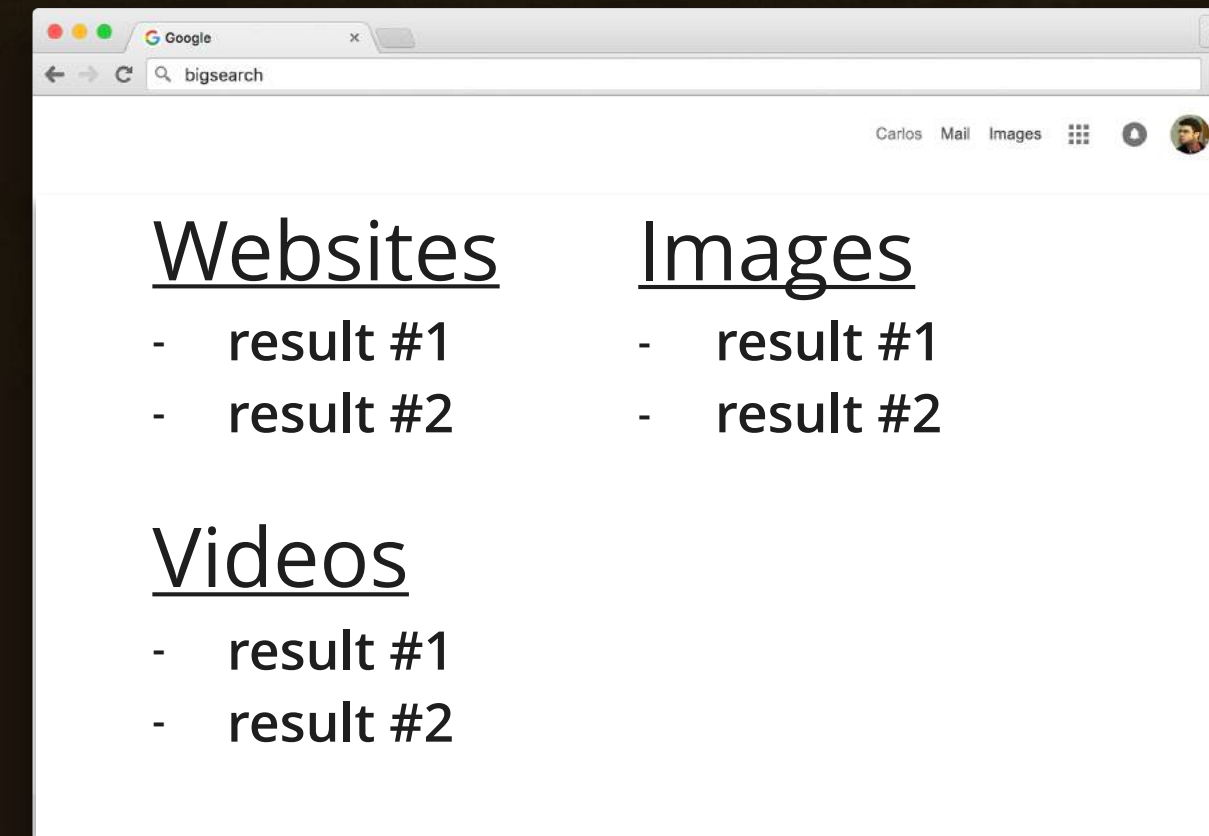


# Parallel Programs

In parallel programs, multiple tasks can be executed **simultaneously** (requires **multi-core** machines).



*Searches websites, images, and videos **at the same time** (really!)*





# Concurrency Allows Parallelism

Concurrency and parallelism are **NOT** the same thing. The former means **independent**, which is a necessary step toward the latter, which means **simultaneous**.

*In short...*

Concurrent *means* Independent



Parallel *means* Simultaneous




Go's **concurrency** model and goroutines make it simple to build **parallel programs** that take advantage of machines with **multiple processors** (most machines today).





# Concurrency With goroutines

A goroutine is a special function that executes **concurrently** with other functions. We create them with the `go` keyword (yes, `go` is also a keyword!).



*On a single-core machine, concurrent code is unlikely to perform better than sequential code.*



# Looping and Printing Names

Let's write a new program that iterates through a slice and invokes a function `printName()` for each item.

```
package main

import "fmt"

func main() {
    names := []string{"Phil", "Noodles", "Barbaro"}
    for _, name := range names {
        printName(name)
    }
}

func printName(n string) {
}
```



# Tracking Duration With the time Tool

The `printName()` function takes one argument and simply prints it to the console. We'll run our program with `go run` and use the Unix `time` command to **track the duration of the execution**.

...

```
func main() {  
    names := []string{"Phil", "Noodles", "Barbaro"}  
    for _, name := range names {  
        printName(name)  
    }  
}
```

```
func printName(n string) {  
    fmt.Println("Name: ", n)  
}
```

*Prints argument to the console*

\$ `time go run main.go`

*Determines duration of command passed to it.*

→  
Name: Phil  
Name: Noodles  
Name: Barbaro  
  
real 0m0.321s

*Less than half a second*



# Heavy Processing Comes Into Play

Let's simulate a **time-consuming task** on `printName()`. We'll do this by adding a very costly mathematical operation using the `math` package from Go's standard library.

```
import (  
    "fmt"  
    "math" ← Import new package.  
)
```

```
func main() { ... }
```

```
func printName(n string) {  
    result := 0.0  
    for i := 0; i < 100000000; i++ {  
        result += math.Pi * math.Sin(float64(len(n)))  
    }  
    fmt.Println("Name: ", n)  
}
```

*Time-consuming computation  
keeps the processor busy!*



# Sequential Tasks Are Blocking

When we add **heavy processing** to `printName()`, we can see a big **time increase on execution time**.

```
...  
func main() {  
    names := []string{"Phil", "Noodles", "Barbaro"}  
    for _, name := range names {  
        printName(name)  
    }  
}
```

*Each call to this function blocks the processor for almost 5 seconds!*

```
func printName(n string) {  
    result := 0.0  
    for i := 0; i < 100000000; i++ {  
        result += math.Pi * math.Sin(float64(len(n)))  
    }  
    fmt.Println("Name: ", n)  
}
```

```
$ time go run main.go
```

```
Name: Phil  
Name: Noodles  
Name: Barbaro  
  
real 0m11.603s
```

*Went from 0.3 to 11.6 seconds!*

*Running sequentially*


```
printName(...) printName(...) printName(...)
```



# Going Concurrent

Go programs **are NOT automatically aware** of newly created goroutines, so the main function **exits before the goroutines are finished**.

```
...  
func main() {  
    names := []string{"Phil", "Noodles", "Barbaro"}  
    for _, name := range names {  
        go printName(name)  
    }  
}  
  
func printName(n string) {  
    ...  
}
```



*Executes each function call on a new goroutine.*

```
$ time go run main.go
```

→ real 0m0.314s

*Back to being fast, but no names listed*

Worry not, my friend. There's a built-in solution for this...



# Adding Synchronization With WaitGroup

On the `sync` package from Go's standard library, there's a `WaitGroup` data type. We can use this type to make our program wait for goroutines to finish.

```
...
import (
    "fmt"
    "sync" ← Import new package.
    "math"
)

func main() {
    names := []string{"Phil", "Noodles", "Barbaro"}
    var wg sync.WaitGroup ← Declare a new variable of
                           the sync.WaitGroup data type.
    ...
}
...
```



# Waiting on goroutines

The `Add` method sets the number of goroutines to wait for, and the `Wait` method prevents the program from exiting before all goroutines being tracked by our `WaitGroup` are finished executing.

...

```
func main() {  
    names := []string{"Phil", "Noodles", "Barbaro"}  
    var wg sync.WaitGroup  
    wg.Add(len(names))  
    for _, name := range names {  
        go printName(name)  
    }  
    wg.Wait()  
}
```

...

*The call to `len()` returns the total number of names...*

*...which is equal to the number of goroutines we create inside the loop.*

*Prevents program from exiting.*



# Updating WaitGroup

The `Done` method must be called **from each function that runs on a** goroutine **once it's finished**. This gives the `WaitGroup` an update — like saying, *"Hey, there's one less goroutine **you need to wait for**."*

```
...  
func main() {  
    var wg sync.WaitGroup  
    wg.Add(len(names))  
  
    ...  
    for _, name := range names {  
        go printName(name, &wg)  
    }  
    wg.Wait()  
}
```

*Must pass a reference to `WaitGroup` so that we call `Done` on the original value and NOT on a copy.*

```
func printName(n string, wg *sync.WaitGroup) {  
  
    ...  
    wg.Done()  
}
```

*Inform the `WaitGroup` that the goroutine running this function is now finished!*



# Single CPU — Concurrent and Synchronized

If we run our final code **specifying a single processor**, there's **no noticeable performance improvement**.

```
...  
func main() {  
    ...  
    var wg sync.WaitGroup  
    wg.Add(len(names))  
    for _, name := range names {  
        go printName(name, &wg)  
    }  
    wg.Wait()  
}
```

```
func printName(n string) {  
    ...  
    wg.Done()  
}
```

```
$ time GOMAXPROCS=1 go run main.go
```

*Run program on a single processor.*

Name: Phil  
Name: Noodles  
Name: Barbaro

real 0m11.675s

*Still slow*

*Running concurrent*

printName(...) printName(...) printName(...) printName(...) printName(...) printName(...)



# Multiple CPUs — Parallel and Synchronized

The Go runtime **defaults to using all processors available**. Most machines today have more than one processor and our concurrent Go code can run **in parallel** with no changes!

*Absence of GOMAXPROCS means all processors available will be used!*

```
...  
func main() {  
    ...  
    var wg sync.WaitGroup  
    wg.Add(len(names))  
    for _, name := range names {  
        go printName(name, &wg)  
    }  
    wg.Wait()  
}
```

```
func printName(n string) {  
    ...  
    wg.Done()  
}
```

*Running in parallel!*

```
$ time go run main.go
```

→ Name: Phil  
Name: Noodles  
Name: Barbaro  
  
real 0m4.172s

*From 11 to 4.1 seconds!*



printName("Phil", &wg)

printName("Noodles", &wg)

printName("Barbaro", &wg)